


PH6232: Machine Learning for Physics applications



CNNs


When inputs have specific relations to each other, or when input information has structural inter-relations which *matter*

For example, images.

CNNs

When inputs have specific relations to each other, or when input information has structural inter-relations which *matter*

Say I was identifying handwritten numerals...



Inputs could be pixel by pixel intensity of the image... is a given pixel black or not.


This is a 200x200 pixel image, so total 40000 numbers.

CNNs


When inputs have specific relations to each other, or when input information has structural inter-relations which *matter*

Say I was identifying handwritten numerals...

But the order of these 40000 numbers carries information!




not same as



Moreover, are all 40k numbers really needed?

Enter convolutions...



Link: Sumit Saha

1 <small>x1</small>	1 <small>x0</small>	1 <small>x1</small>	0	0
0 <small>x0</small>	1 <small>x1</small>	1 <small>x0</small>	1	0
0 <small>x1</small>	0 <small>x0</small>	1 <small>x1</small>	1	1
0	0	1	1	0
0	1	1	0	0

Image

4		

Convolved
Feature


Here, the filter (or kernel) is run over the image moving from left to right, top to bottom

Filter is


1	0	1
0	1	0
1	0	1

, performing a Hadamard product.


Terminology



3x3 filter, stride=1




3x3 filter, stride=2



3x3 filter, stride=2, padded


Stride is step (in x or y) between successive operations

Terminology




3x3 filter, stride=1

4x4 \rightarrow 2x2



3x3 filter, stride=2

5x5 \rightarrow 2x2



3x3 filter, stride=2, padded


5x5 is padded to 7x7

5x5 \rightarrow 3x3

Input $n \times n$, output is $m \times m$


$m = (\text{width} - \text{kernel}) / \text{stride} + 1$ (round up)

Terminology




3x3 filter, stride=1

4x4 \rightarrow 2x2



3x3 filter, stride=2

5x5 \rightarrow 2x2



3x3 filter, stride=2, padded

5x5 is padded to 7x7


5x5 \rightarrow 3x3

Input $n \times n$, output is $m \times m$

$m = (\text{width} - \text{kernel}) / \text{stride} + 1$ (round up)


Padding ensures edge of image gets
read in enough too.

MaxPooling




For 2x2 maxpooling, take the max value of a 2x2 block, and stride across


For 2x2 averagepooling, take the average value of 2x2 block, and stride across




Putting it together



Flatten



Let's take an example





Here we have 16 detectors,
yielding 16 points.

Charged particles bend in a magnetic field. The magnitude of their momentum decides the radius of curvature.

We measure the “points” where the particle intersects our detector. A “fit” to the points will give the trajectory of the particle.





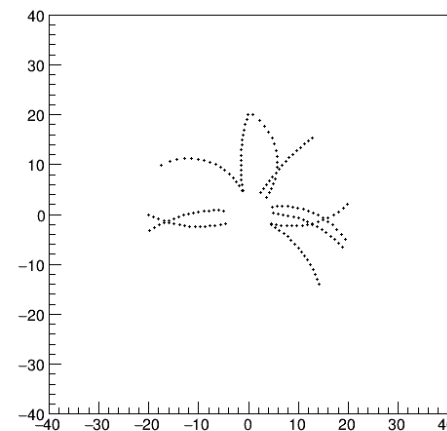



Figure shows one charged particle trajectory in a “spectrometer” designed to measure momentum.

Let's take an example



Right figure shows three charged particle trajectories. One of these is “straighter” indicating a larger momentum (large radius of curvature = large momentum)

Each image has 10 charged particles, one of which may be of high momentum. Can you tell which of these contain a “high momentum” particle?



We setup a CNN

```
model = Sequential()
model.add(Conv2D(32, (5,5), strides=(1,1),
activation='relu',kernel_initializer='he_uniform', padding='SAME',
input_shape=(IMG_HEIGHT,IMG_WIDTH,1)))
model.add(MaxPool2D((2,2), padding='SAME'))
model.add(BatchNormalization())
model.add(Conv2D(64, (3,3), strides=(1,1),
activation='relu',kernel_initializer='he_uniform', padding='SAME'))
model.add(MaxPool2D((2,2), padding='SAME'))
model.add(BatchNormalization())
model.add(Flatten())
model.add(Dense(64, activation='relu', kernel_initializer='he_uniform'))
model.add(Dense(32, activation='relu', kernel_initializer='he_uniform'))
model.add(Dense(1, activation='sigmoid'))



model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
history = model.fit(train_data, epochs=10, verbose=0, validation_data=test_data,
callbacks=cb)
```

Runtime is slow...



- 1) Resize images from 512x512 to 128x128 to save computation
- 2) Save model for best val_accuracy while training

```
Epoch 1: val_accuracy improved from -inf to 0.50100, saving model to best_model.h5
Epoch 2: val_accuracy improved from 0.50100 to 0.50640, saving model to best_model.h5
Epoch 3: val_accuracy did not improve from 0.50640
Epoch 4: val_accuracy improved from 0.50640 to 0.68570, saving model to best_model.h5
Epoch 5: val_accuracy did not improve from 0.68570
Epoch 6: val_accuracy did not improve from 0.68570
Epoch 7: val_accuracy did not improve from 0.68570
Epoch 8: val_accuracy improved from 0.68570 to 0.76360, saving model to best_model.h5
Epoch 9: val_accuracy improved from 0.76360 to 0.79970, saving model to best_model.h5
Epoch 10: val_accuracy improved from 0.79970 to 0.84110, saving model to best_model.h5
```

Results



Results



Code etc.

Not expecting you to implement this code.

But, the code (as a python file) and the input data is linked from the course webpage.

If you feel like playing with the CNN, go ahead. Just remember, its pretty resource-intensive.